


A-Eye App

For Image Captioning

AC 215: Advanced Practical Data Science, MLOps
Group Name: Pinkdrink
Heather Liu, Anita Mahinpei





- Background
- Proposed Solution
- Project Workflow
- Data
- Models
- App Architecture
- Deployment
- Conclusions
- References



- In 2020, approximately 43.3 million people were blind and 295 million people had moderate to severe vision impairment worldwide (Bourne et al 2021).
- Many people with visual impairments rely on screen readers in order to access the internet through audio and thus depend on image captions (Yesilada et al 2004).
- Studies have revealed that most websites are not accessible to visually impaired individuals due to missing image captions/alt text (McEwan et al 2007).



Therefore, with the ever-increasing prevalence of online services, accurate image caption is an important priority for improving accessibility.



A-Eye App for Image Captioning



“ A man in the surf ”

“ A man in surfing ”

“ A man on a surfboard ”



User Input Image

Image Caption
Translation

Read Out



To tackle this problem, we developed a web application that allows users to upload images and have them be captioned by three different machine learning models. We specifically focused on creating captions for images of objects and scenery. Our app generates and displays three possible captions for a user provided image. Since the intended audiences are visually impaired individuals, our app provides a text-to-audio functionality so that the generated captions can be read out loud to the user if desired. To support a broader range of audiences, we also allow users to select the language they wish to use for the generated captions.



Proof Of Concept (POC)

- Download MSCOCO data
- Perform EDA to verify data
- Set up data pipeline (i.e. resize all images to a fixed size, normalize pixel values, tokenize the captions, store in tf dataset)
- Experiment with some baseline models trained on a subset of the dataset
- Validate captioning results on unseen images

Prototype

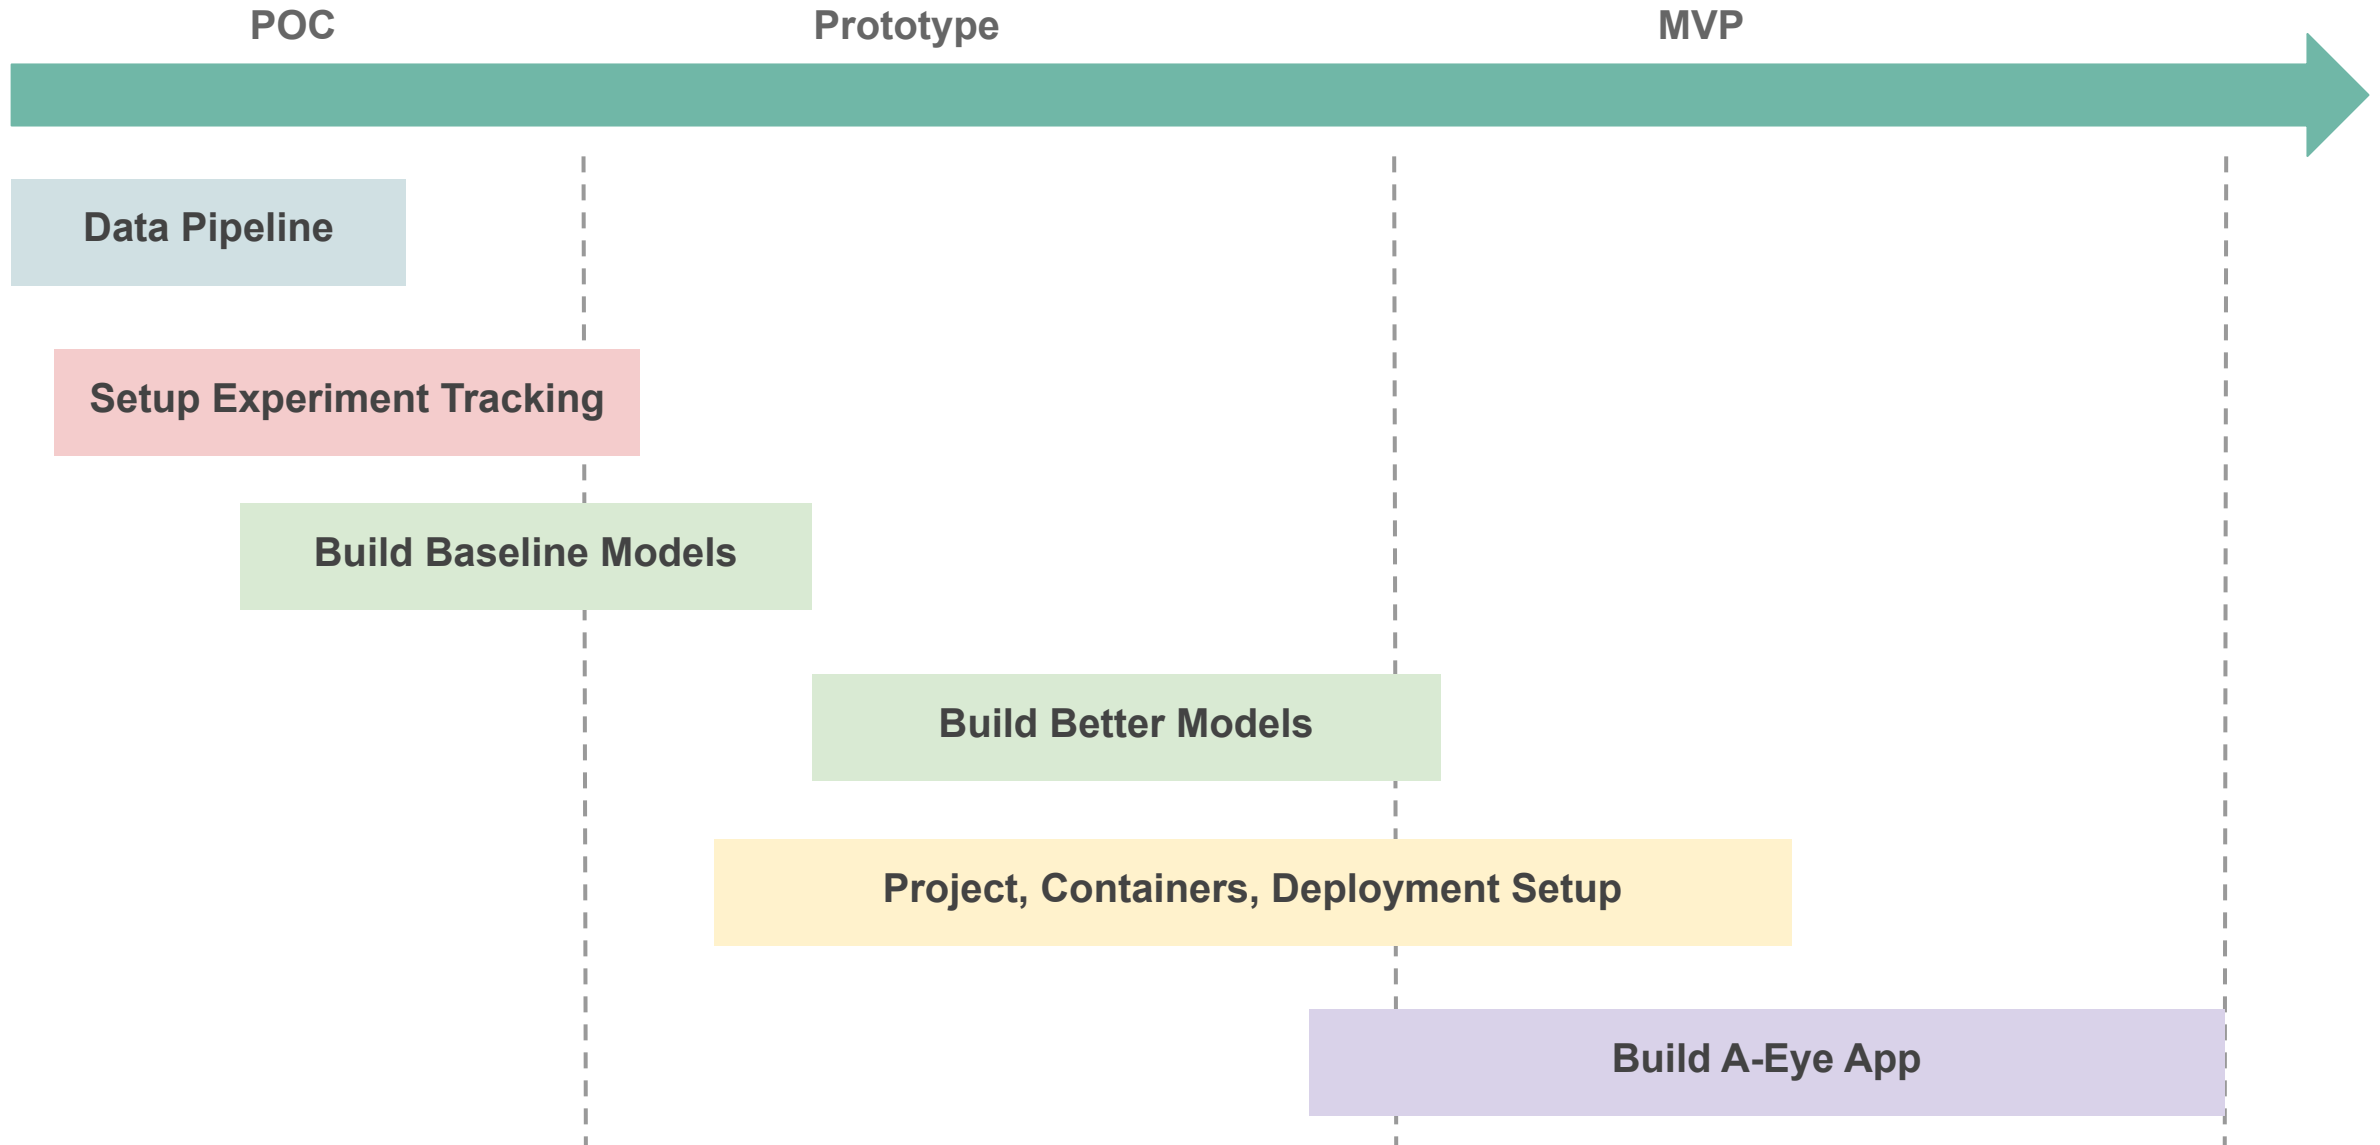
- Create a mockup of screens to see how the app could look like
- Deploy one model to Fast API to service model predictions as an API

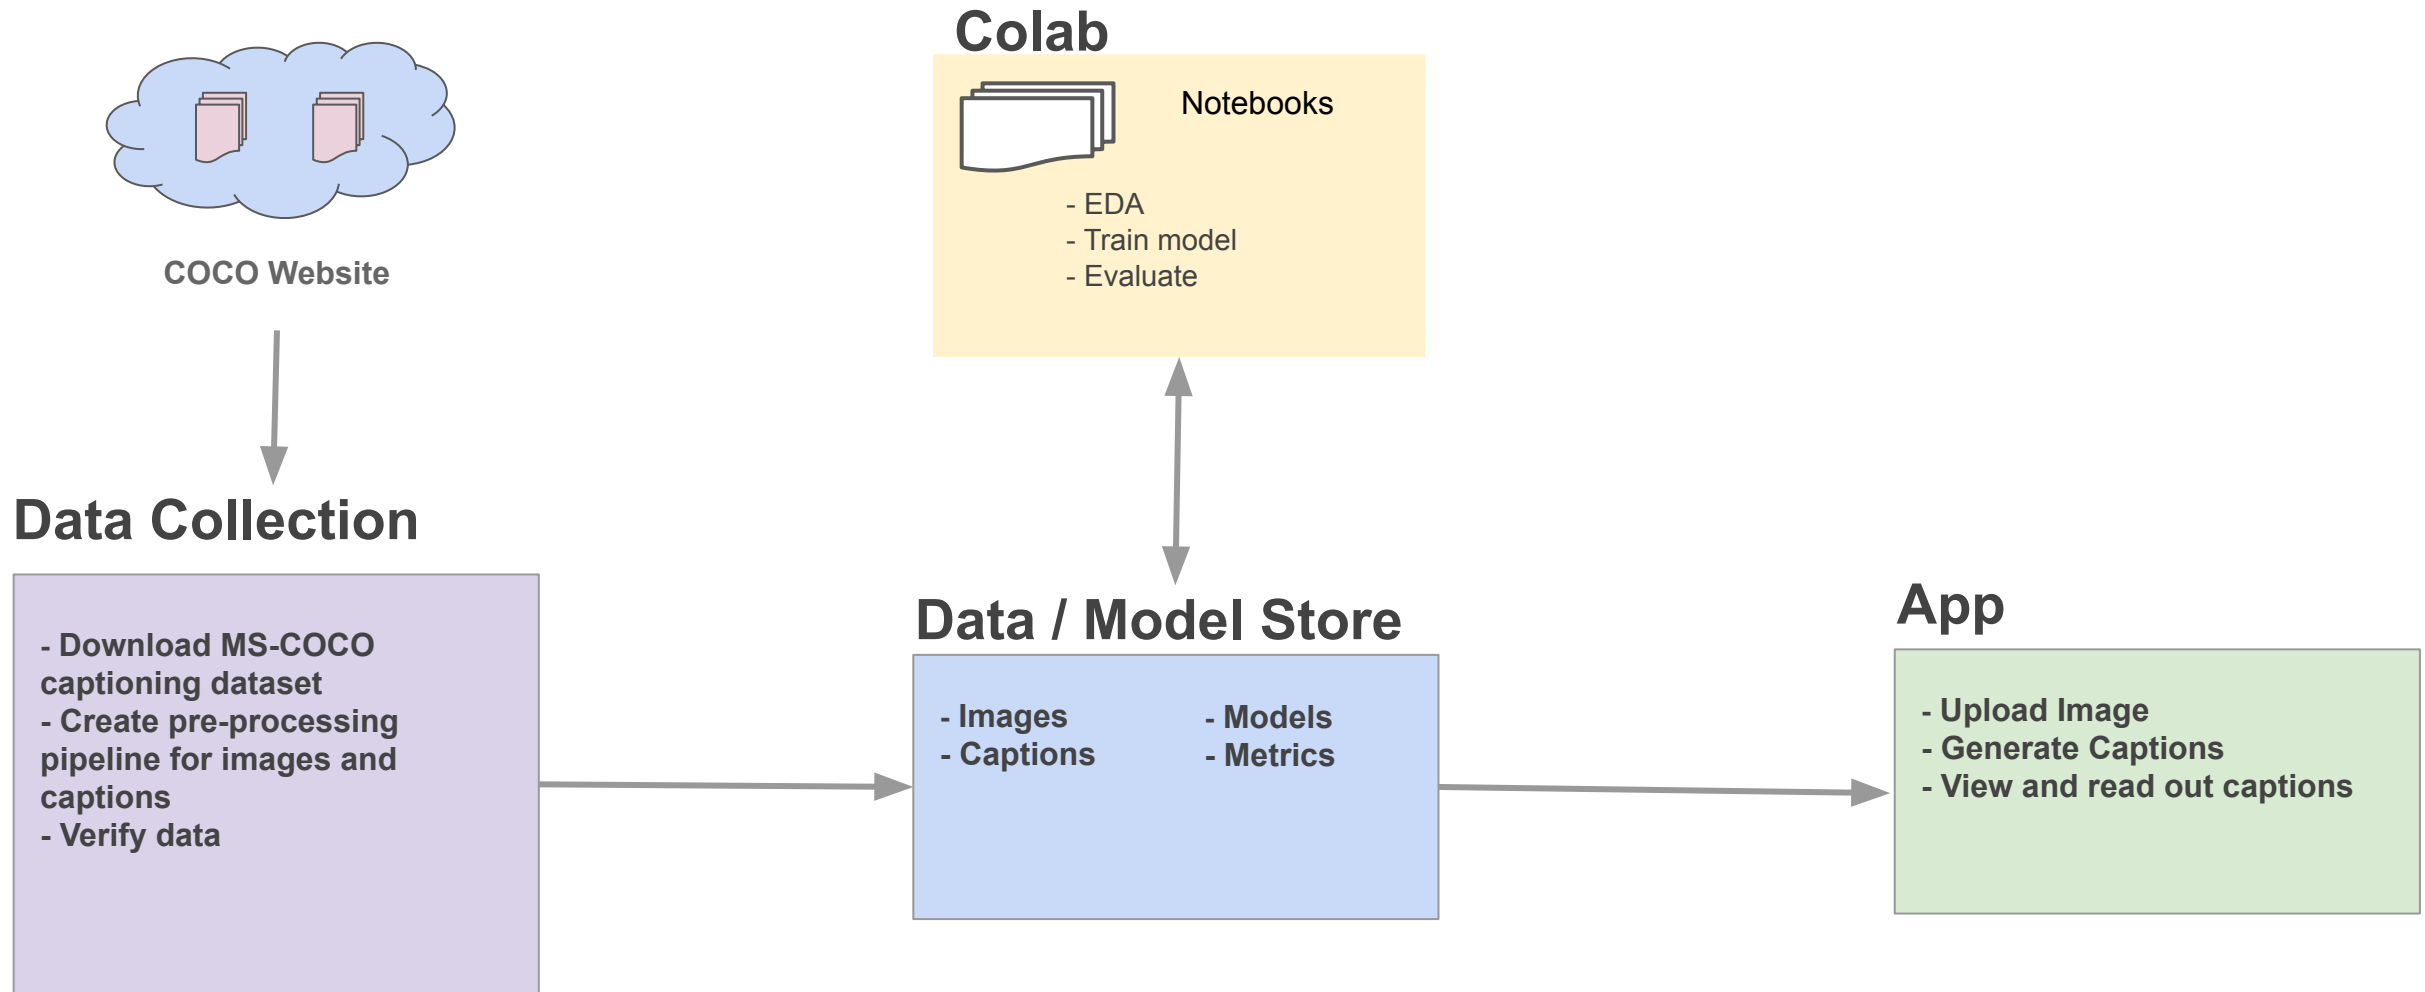
Minimum Viable Product (MVP)

- Create App to caption images
- API Server for uploading images and predicting using best model

Project Workflow

A-Eye App







We used the Microsoft [Common Objects in Context](#) (COCO) data for our project. COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- Object segmentation
- Recognition in context
- Superpixel stuff segmentation
- 330K images (>200K labeled)
- 1.5 million object instances
- 80 object categories
- 91 stuff categories
- 5 to 7 captions per image
- 250,000 people with keypoints



We used the images and the caption labels for training models.


Sample image with 5 gold captions:



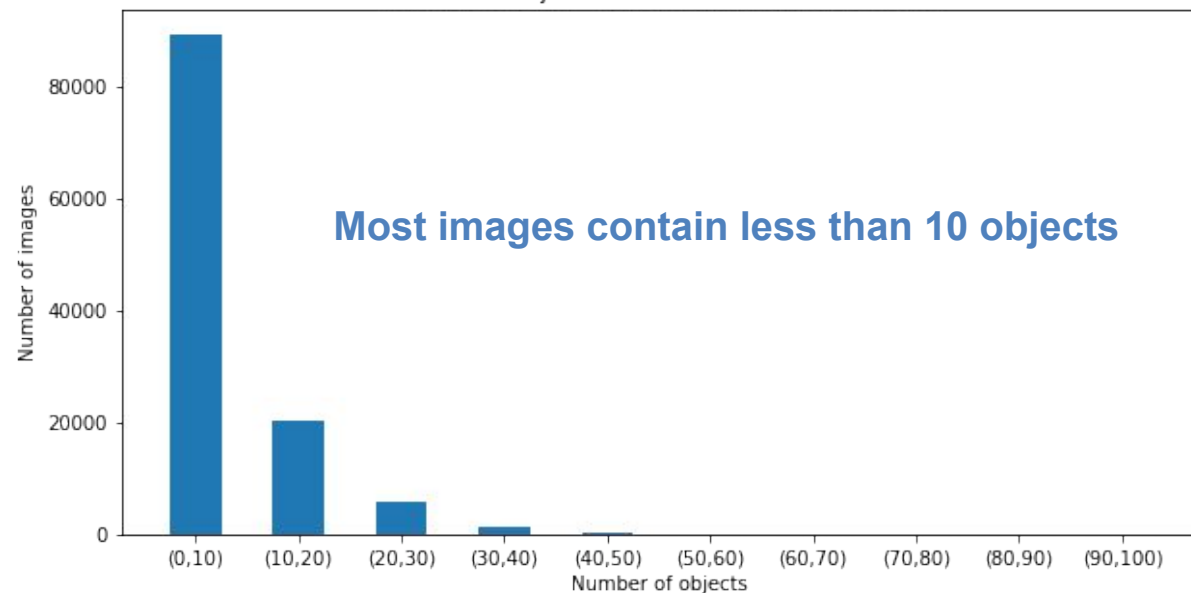
A couple of people riding on top of a wave on surfboards.
A man rides a white surfboard near another person in the ocean.
one person surfing one person laying on a surfboard
The guy is riding the wave as a girl watches.
Surfers surfing in the ocean on a clear day.

Data: EDA Results

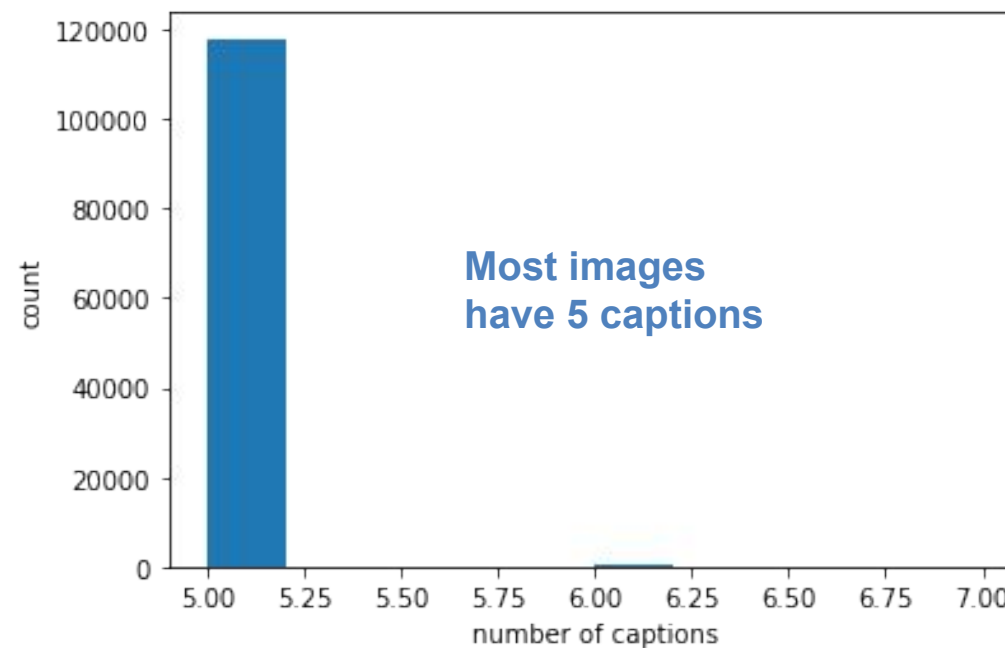
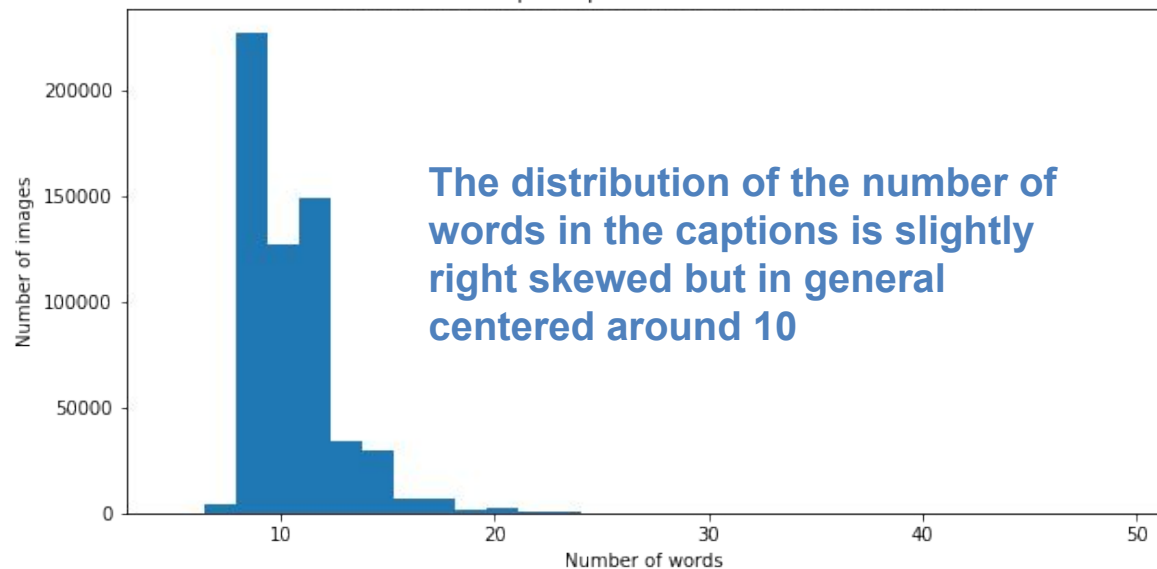
A-Eye App






Number of objects distribution over the dataset



Number of words per caption distribution over the dataset







In terms of deep learning models, both computer vision and natural language processing models will be used.

- **Computer Vision:** Pre-trained, frozen CNN architectures (e.g. VGG, Inception) will be used to extract features from the images. (Note: we decided to use frozen encoders trained on ImageNet because we ran a trial without freezing layers and observed a lot of over-fitting.)
- **Language:** RNN and LSTM structures
- **Attention:** As described in the paper, “[Show, Attend, and Tell](#)”, models that incorporate attention to the image feature map perform, have improved performance. We will use the attention mechanism described in this paper in one of our models.



We have trained the following models:

1. **Inception-GRU:**

- Extract feature map from the last CNN layer of frozen InceptionV3 with ImageNet weights
- Attend to feature map
- Decoder with an embedding layer (not pre-trained), a GRU layer and 2 fully-connected layers

2. **Inception-LSTM:**


- Extract feature map from the last CNN layer of frozen InceptionV3 with ImageNet weights
- Attend to feature map
- Decoder with an embedding_layer(not pre-trained), a LSTM layer and 1 fully-connected layer

3. **VGG-LSTM:**

- Extract feature map from the last CNN layer of frozen VGG16 with ImageNet weights
- Language feature extractor with an embedding_layer(not pre-trained), a dropout layer, and an LSTM layer. (This extracts features from the previous caption words to predict the next word)
- Combine the output of the language and image feature extractors and feed it into a fully-connected layer to predict the next word of the caption



We compared the performance of our models using the loss function and BLEU-4 scores on held-out test data. The BLEU score is a method of comparing generated texts against a set of golden text labels (in our case image captions) as described in [this paper](#).



We trained with 82783 images each with 5 captions. We used 500 unseen image-caption pairs from the COCO dataset to calculate the BLEU-4 score.

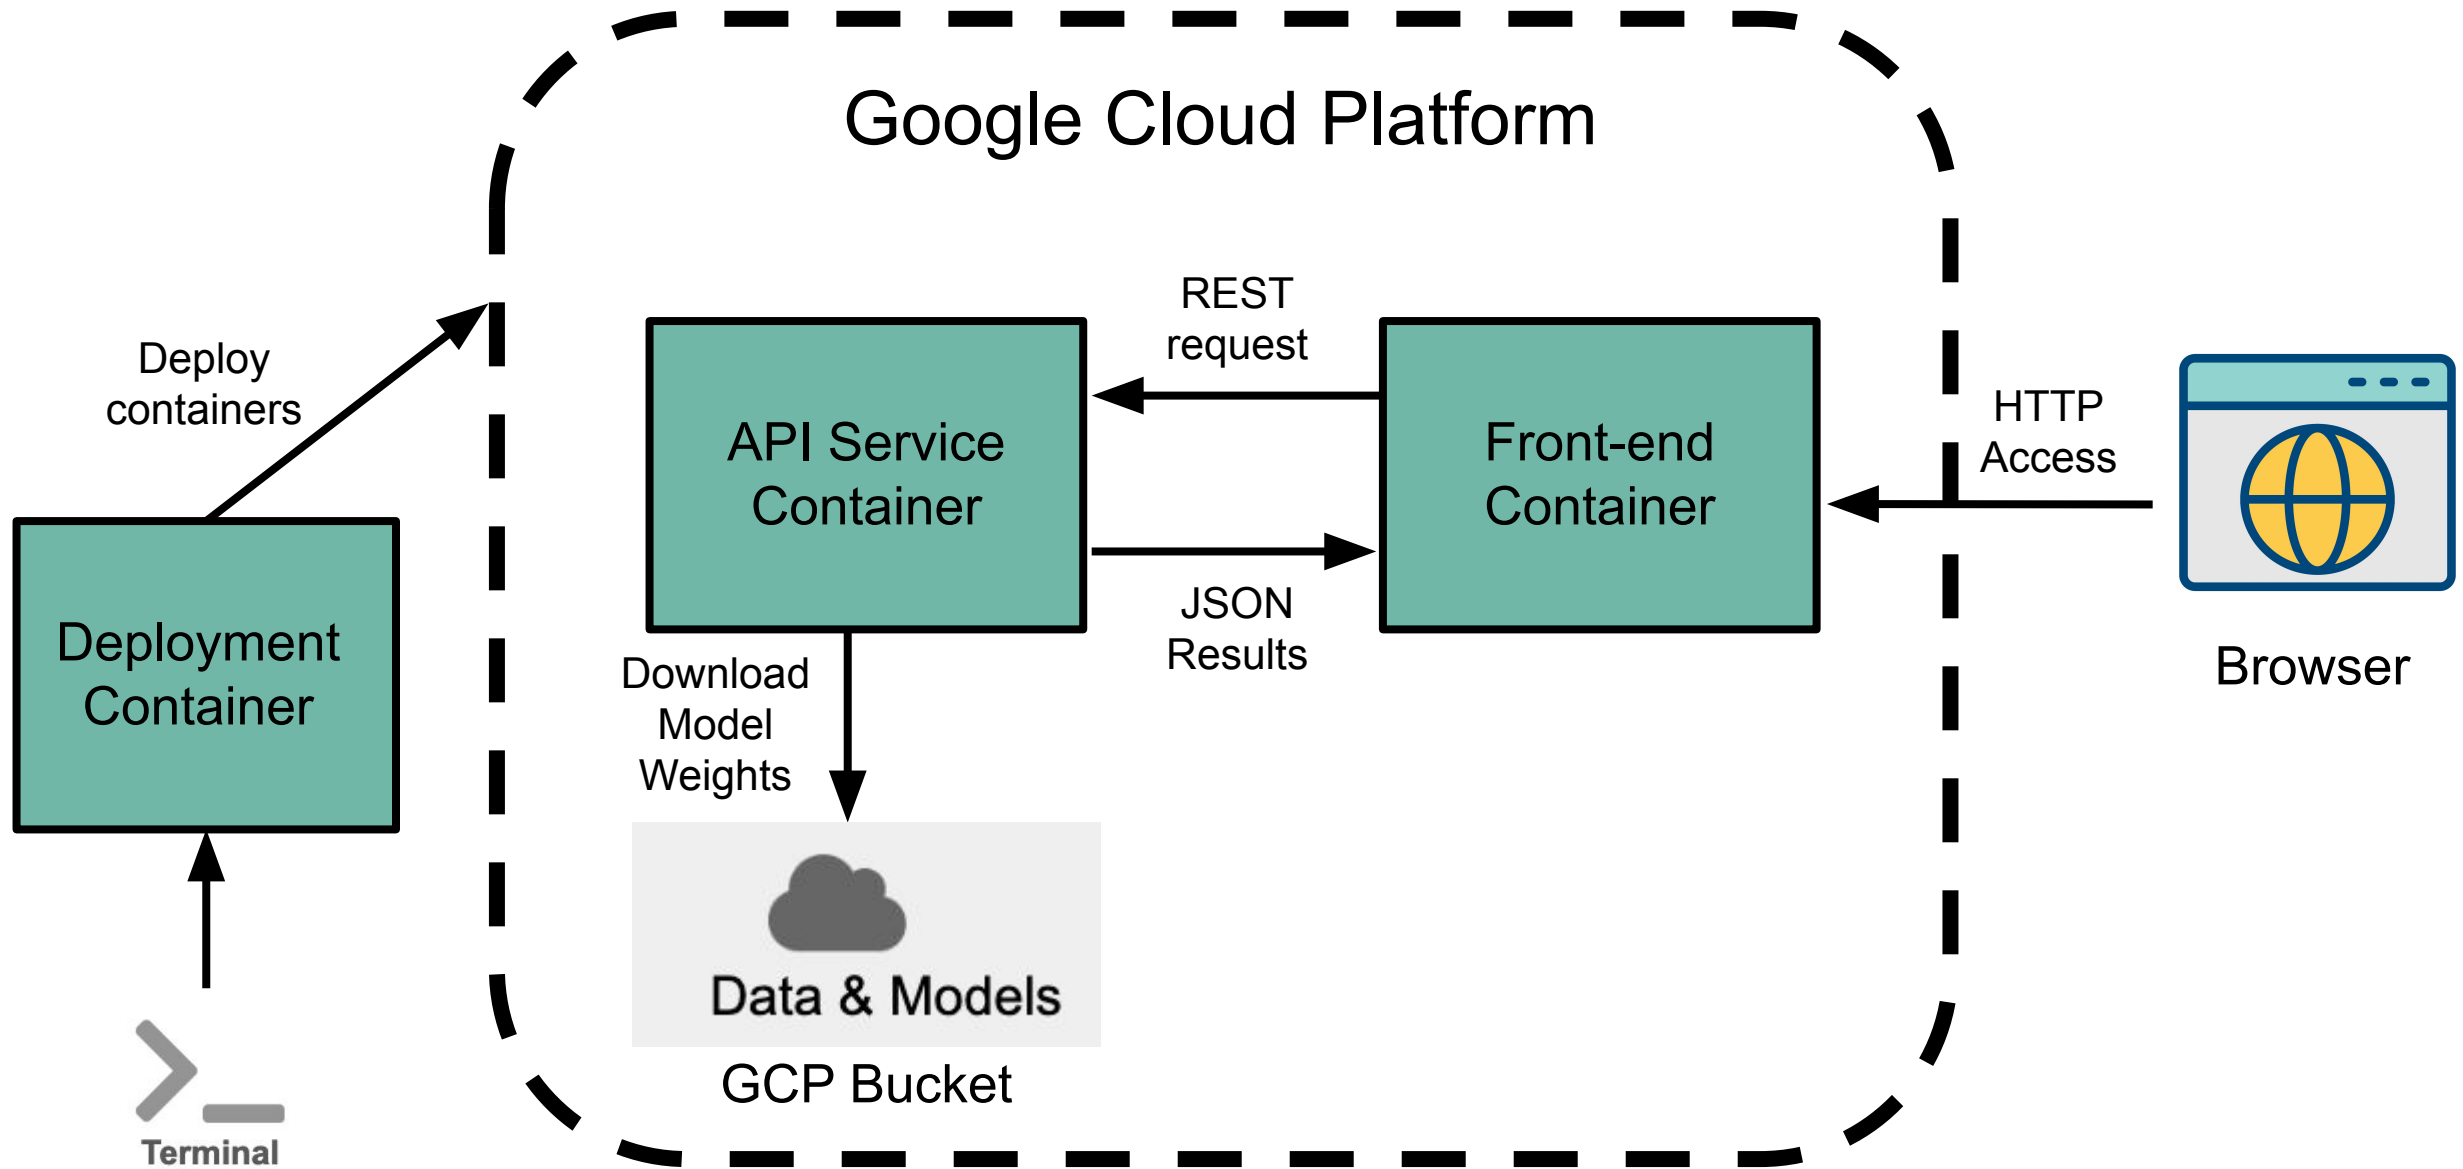
We can see that adding an attention mechanism helps a lot with performance. Both the Inception-GRU and Inception-LSTM models which used attention have a higher BLEU score than the model that didn't use attention. Using an LSTM layer in the decoder also helps improve performance significantly as opposed to using a GRU layer.

Model	Test BLEU-4
Inception-GRU	0.47476
Inception-LSTM	0.60926
VGG-LSTM	0.43329



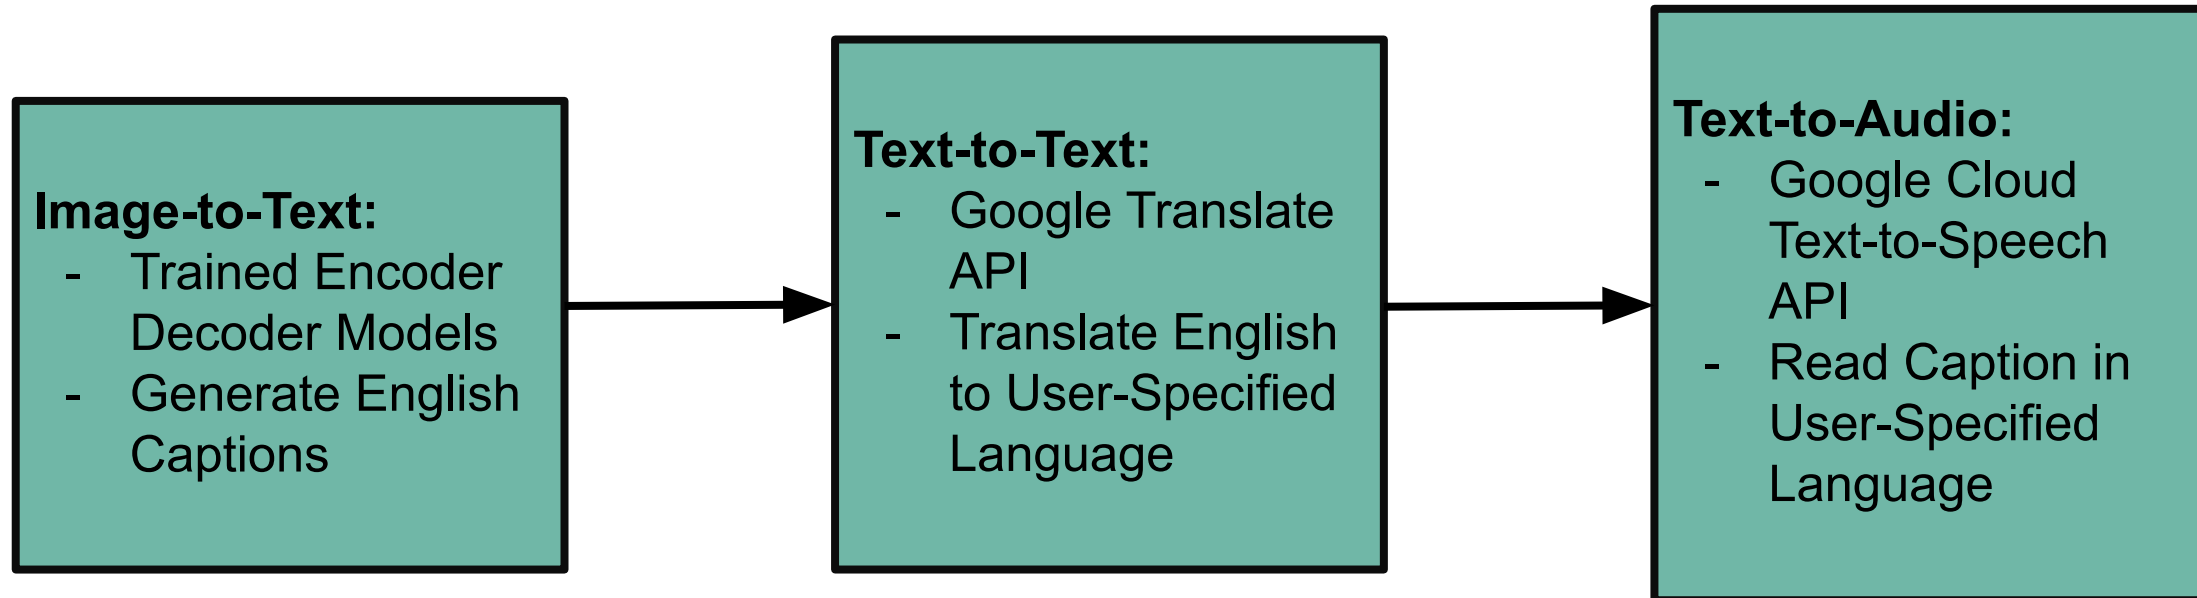
Our app constitutes three docker containers:

1. api-service:
 - a. one endpoint to generate captions and one that returns the caption audios
 - b. written in Python and uses the Google Cloud APIs
2. frontend-react:
 - a. Allows users to upload an image, view the captions and caption audios and change the language settings of the app.
 - b. Written in Javascript and CSS and uses the React library.
3. deployment:
 - a. Used to deploy the app to GCP using Ansible scripts.





API Service Process Flow:





We used Ansible scripts to deploy our app on GCP and Kubernetes.

- Deployment to GCP
 - <http://35.222.164.164/>
- Deployment to K8s Cluster
 - <http://35.202.124.222.sslip.io/>

Note: we have shut them down for now to save GCP credit.

a-eye-app-project

Filter

Enter property name or value

?

Name ↑	Hostname ?	Visibility ?
 a-eye-app-api-service	gcr.io	Private
 a-eye-app-frontend-react	gcr.io	Private

VM instances

CREATE SCHEDULE

DELETE

RESET

OPERATIONS

HELP ASSISTANT

SHOW INFO PANEL

LEARN

INSTANCES

INSTANCE SCHEDULE






VM instances are highly configurable virtual machines for running workloads on Google infrastructure. [Learn more](#)

Filter

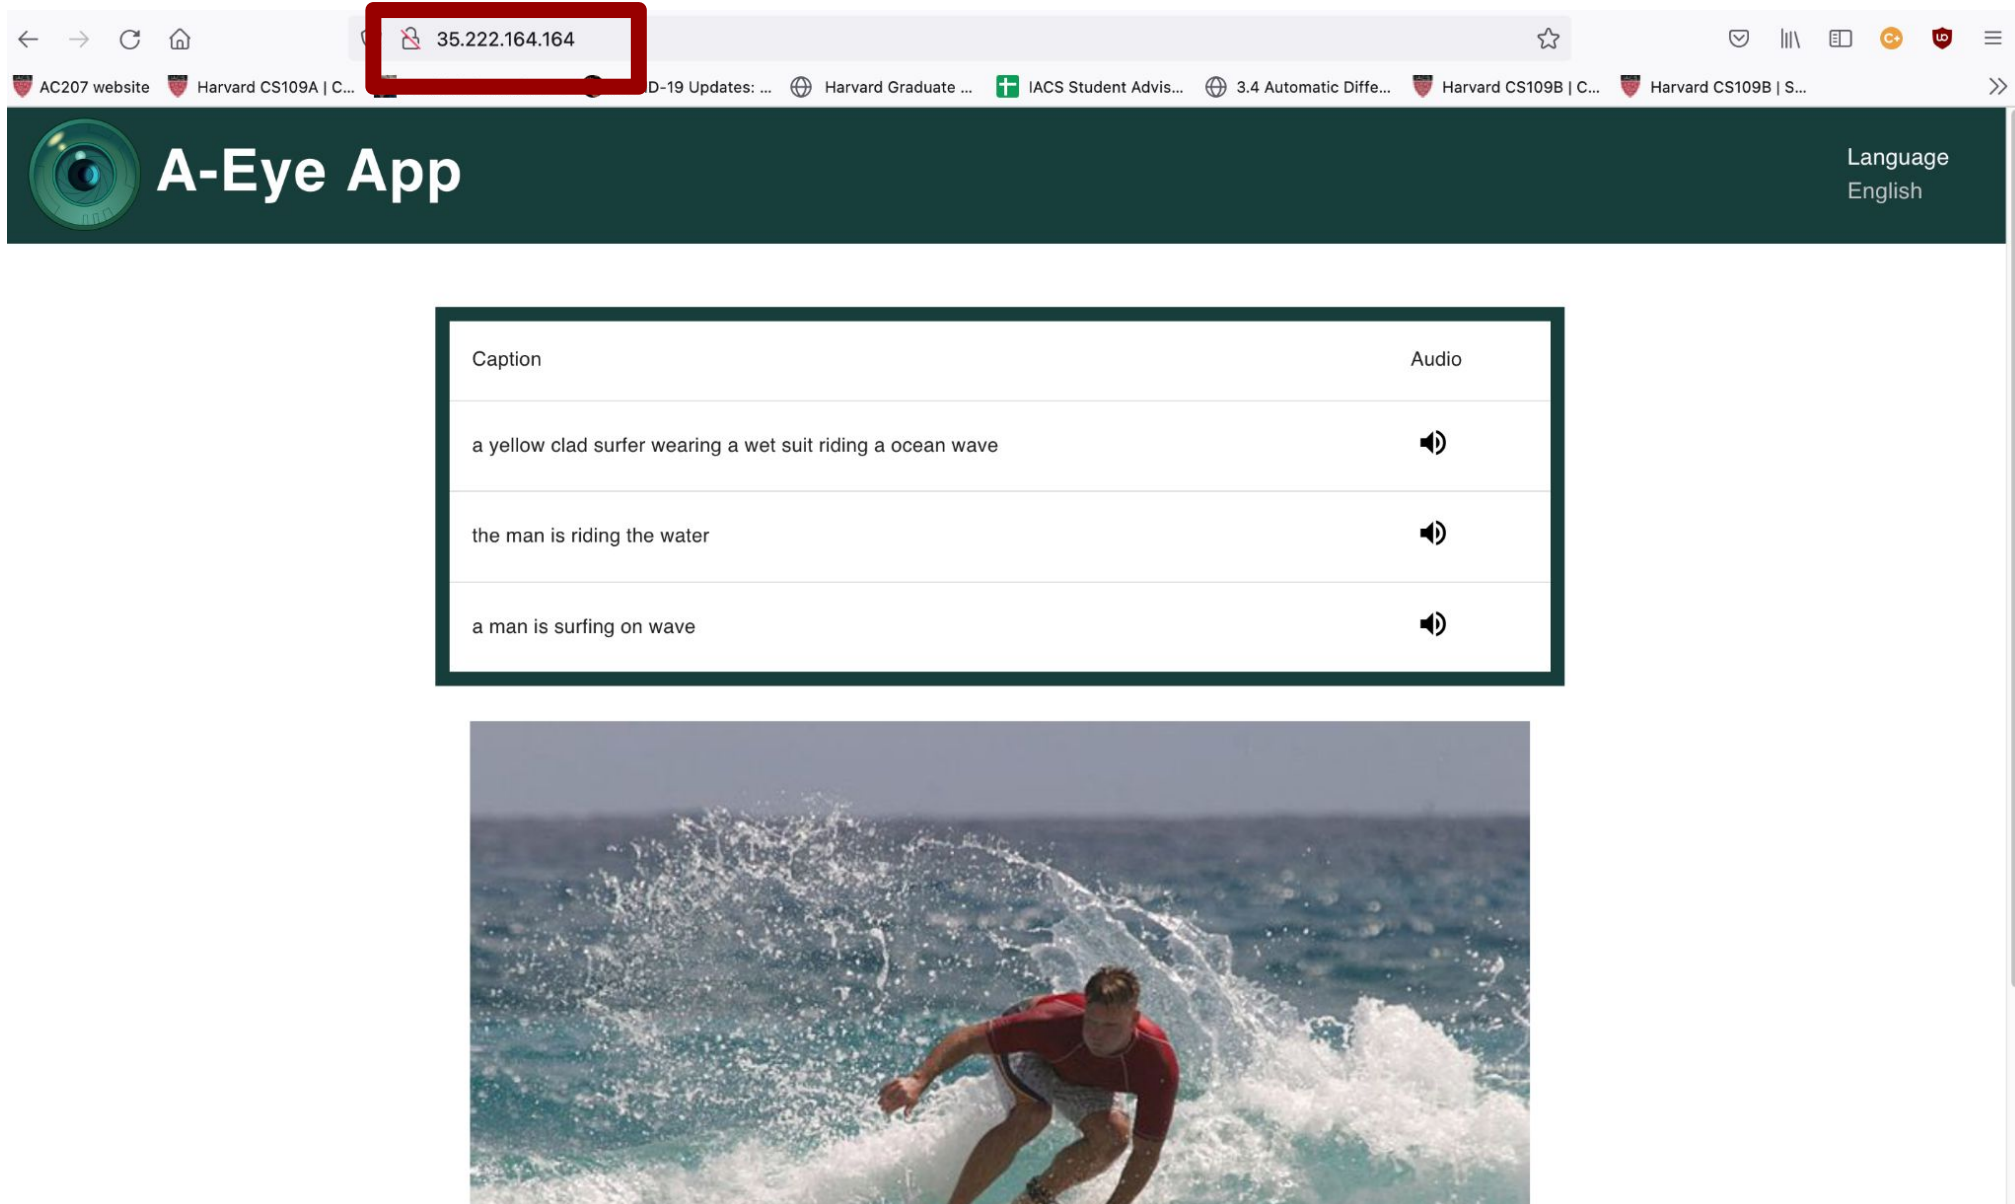
Enter property name or value

?

|||

Status	Name ↑	Zone	Recommendations	In use by	Internal IP	External IP	Connect
	a-eye-app	us-central1-a			10.128.0.4 (nic0)	35.222.164.164  	SSH ▾ ⋮
	gke-a-eye-app-cluster-default-pool-74a0105f-2tf4	us-central1-a		gke-a-eye-app-cluster-default-pool-...	10.128.0.8 (nic0)	35.202.124.222	SSH ▾ ⋮
	gke-a-eye-app-cluster-default-pool-74a0105f-xrvj	us-central1-a		gke-a-eye-app-cluster-default-pool-...	10.128.0.7 (nic0)	35.184.212.54	SSH ▾ ⋮

A-Eye App



Deployment: Kubernetes

Kubernetes clusters

+ CREATE

+ DEPLOY

REFRESH

OPERATIONS

OVERVIEW

COST OPTIMIZATION

PREVIEW

Filter

Enter property name or value

?

⋮

<input type="checkbox"/> Status	Name ↑	Location	Number of nodes	Total vCPUs	Total memory	Notifications	Labels
<input type="checkbox"/>	a-eye-app-cluster	us-central1-a	2	2	7.5 GB	—	⋮

Services & Ingress

REFRESH

+ CREATE INGRESS

DELETE

Cluster

Namespace

RESET

SAVE

SERVICES

INGRESS

Services are sets of Pods with a network endpoint that can be used for discovery and load balancing. Ingresses are collections of rules for routing external HTTP(S) traffic to Services.

Filter

Is system object : False

Filter services and ingresses

×

?

⋮

<input type="checkbox"/> Name ↑	Status	Type	Endpoints	Pods	Namespace	Clusters
<input type="checkbox"/> api	OK	Node Port	10.24.14.220:9000 TCP	1/1	a-eye-app-cluster-namespace	a-eye-app-cluster
<input type="checkbox"/> frontend	OK	Node Port	10.24.12.145:80 TCP	1/1	a-eye-app-cluster-namespace	a-eye-app-cluster
<input type="checkbox"/> nginx-ingress-nginx-ingress	OK	External load balancer	35.202.124.222:80	1/1	a-eye-app-cluster-namespace	a-eye-app-cluster

35.202.124.222.sslip.io

AC207 website

Harvard CS109A | C...

Unofficial Waterloo ...

COVID-19 Updates: ...


Harvard Graduate ...

IACS Student Advis...

3.4 Automatic Diffe...




Harvard CS109B | C...


Harvard CS109B | S...



A-Eye App




Language
Chinese

Caption	Audio
一个男人正在摆动网球拍和球拍	
两个人在草地上玩耍	
拿着在网球场的人网球球拍	







App Demo

Caption	Audio
several people that are playing frisbee on a field	
two children are playing with ball	
two children are playing soccer on field	





Caption	Audio
Une adolescente portant sur le bord du paddle dans l'océan	
Jeune garçon en chemise rouge joue dans l'eau	
Deux enfants jouent dans de l'eau avec leurs planches de surf	





Caption	Audio
冲浪板上的一个男人在海洋中穿着黑色潜水服的冲浪板	
那个男人骑着水	
一个男人在波浪上冲浪	





Caption	Audio
Una pequeña pizza sentada en la parte superior de una sartén.	
Dos chicas están jugando en la hierba.	
Pizza con queso y verduras en él.	





- We created a simple web application that automatically captions images in four different languages (English, French, Spanish, Chinese).
- Future extensions can:
 - Support more languages
 - Combine multiple image captioning datasets (MS-COCO, Flickr30k, etc) to train better models with more data
 - Train on and support other types of images that are commonly found on websites such as infographics
 - Use pre-trained language embeddings (e.g. GloVe or Word2Vec) in the decoder



- Github Repo: https://github.com/pinkpinkdrink/AC215_pinkdrink
- Medium Post Draft:
<https://medium.com/@amahinpei/a-eye-image-captioning-app-3bf7c1d11e91>
- Video Recording: <https://youtu.be/1GRi85gAUbw>



1. Alkhathlan, Mallak, et al. “‘Honestly I Never Really Thought about Adding a Description’: Why Highly Engaged Tweets Are Inaccessible.” *Human-Computer Interaction—INTERACT 2021*, 2021, pp. 373–395., doi:10.1007/978-3-030-85623-6_23.
2. Bourne, Rupert, et al. “Trends in Prevalence of Blindness and Distance and near Vision Impairment over 30 Years: An Analysis for the Global Burden of Disease Study.” *The Lancet Global Health*, vol. 9, no. 2, 2021, doi:10.1016/s2214-109x(20)30425-3.
3. Brownlee, Jason. “How to Develop a Deep Learning Photo Caption Generator from Scratch.” *Machine Learning Mastery: How to Develop a Deep Learning Photo Caption Generator from Scratch*, 23 Dec. 2020, machinelearningmastery.com/develop-a-deep-learning-caption-generation-model-in-python/.
4. Lin, Tsung-Yi, et al. “Microsoft Coco: Common Objects in Context.” *Computer Vision—ECCV 2014*, 2014, pp. 740–755., doi:10.1007/978-3-319-10602-1_48.
5. McEwan, Tom, and Ben Weerts. “Alt Text and Basic Accessibility.” *Electronic Workshops in Computing*, 2007, doi:10.14236/ewic/hci2007.64.
6. Papineni, Kishore, et al. “Bleu.” *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics—ACL ’02*, 2001, doi:10.3115/1073083.1073135.
7. Sharma, Piyush, et al. “Conceptual Captions: A Cleaned, Hypernymed, Image Alt-Text Dataset for Automatic Image Captioning.” *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2018, doi:10.18653/v1/p18-1238.
8. Xu, Kelvin, et al. “Show, Attend and Tell: Neural Image Caption Generation with Visual Attention.” *International Conference on Machine Learning*, 2015.
9. Yesilada, Yeliz, et al. “Screen Readers Cannot See.” *Lecture Notes in Computer Science*, 2004, pp. 445–458., doi:10.1007/978-3-540-27834-4_55.